

# Planning in deep reinforcement learning

Alexander Lyzhov

MDP

Planning

Planning in deep RL

MuZero architecture

MuZero results

# MDP

Trajectory:  $s_0, a_1, r_1, s_1, a_2, r_2, s_2, a_3, r_3, s_3, \dots$

Dynamics:  $\mathbb{P}(s', r \mid s, a)$

Model:  $\pi(a \mid s)$

Return:  $\underline{g_t} = r_{t+1} + \underbrace{\gamma r_{t+2} + \gamma^2 r_{t+3} + \gamma^3 r_{t+4} + \dots}_{< 1}$

Value function:  $v_\pi(s) = \mathbb{E}_\pi(g_t \mid s_t = s)$

Q-value function:  $q_\pi(s, a) = \mathbb{E}_\pi(g_t \mid s_t = s, a_t = a)$

$$g_t = r_{t+1} + \gamma g_{t+1}$$

$$v_\pi(s) = \mathbb{E}_{a_{t+1}, r_{t+1}, s_{t+1}}(r_{t+1} + \gamma v_\pi(s_{t+1}) \mid s_t = s)$$

$$q_\pi(s, a) = \mathbb{E}_{r_{t+1}, s_{t+1}, a_{t+2}}(r_{t+1} + \gamma q_\pi(s_{t+1}, a_{t+2}) \mid s_t = s, a_{t+1} = a)$$

# Planning

Environment model - anything that can be used to predict how the env responds to actions.

Planning - anything that improves policy with a model.

- ▶ Background planning - to improve the policy when training
- ▶ Online planning - to make better decision with trained policy

Model-based vs model-free - continuum.

Analogies: self-supervision, augmentation, actor-critic methods, 2-player game [1].

## Planning: pros and cons

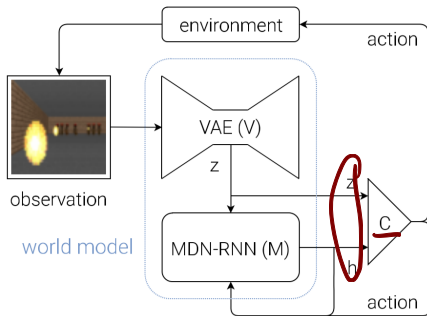


- + works in domains with complex dynamics (e.g. Go)
- + more sample-efficient
- + better transfer, including zero-shot adaptation to new tasks [2]
- + better offline RL [3]
- takes time/memory to build environment model and use it
- more complex architecture, more hyperparameters


plan 2 explore  
 $p: s, a \rightarrow r$

# Model-based deep RL in discrete environments

World models[4]



SimPLe [5]




Also: Value Iteration Networks [6], Value Prediction Networks [7], Predictron [8].  
But SOTA was model-free before MuZero (R2D2 on Atari).

MRP


# MuZero[9]

Discrete state & action spaces, deterministic dynamics: Atari, Go, Chess, Shogi.  
Idea: only predict those aspects of the future that help to pick the best action.




Want to predict:

$$\begin{array}{l} \text{policy} \\ \text{value} \\ \text{reward} \end{array} \left\{ \begin{array}{l} \pi(a_{t+k+1} | o_{1:t}, a_{1:t+k}) \\ E(r_{t+k+1} + \gamma r_{t+k+2} + \dots | o_{1:t}, a_{1:t+k}) \\ p(r_{t+k} | o_{1:t}, a_{1:t+k}) \end{array} \right. \quad \forall k$$



- 1) novelty  
2) value  
3) policy

$h$  representation  
 $f$  prediction  
 $g$  dynamics




MuZero planning: MCTS

- 1) Repeatedly select actions by balancing exploration with exploitation on each timestep
- 2) Expand tree with leaf state, remember final reward and transition
- 3) Update values  $Q(s, a)$  on trajectory using final value and intermediate rewards


Result: visit counts  $N(s_0, a)$

# MuZero acting





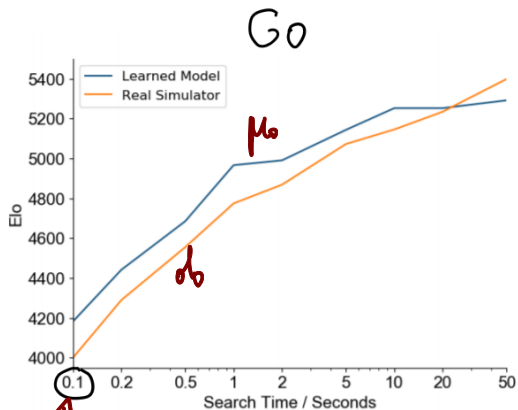
# MuZero training



# Hyperparameters

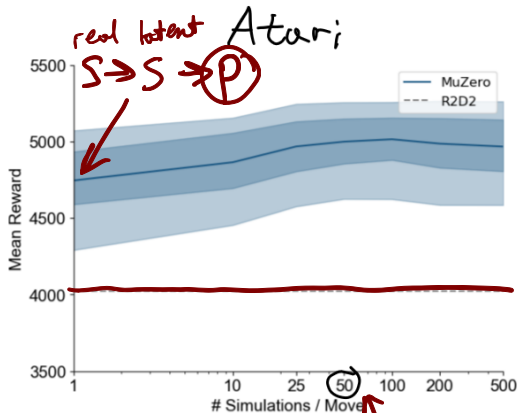
- ▶ Deep resnets <sup>16 blocks</sup> for dynamics  $g$  and representation  $h$ , shallow convolutional net with FC for prediction module  $f$
- ▶ Planning horizon  $K = 5$  — layers in tree for backprop
- ▶ Bootstrapping values: 10 steps for Atari, to end of game for board games
- ▶ 1M mini-batches of size 1-2K. Prioritized replay for Atari, uniform replay for board games

# MuZero results



800 rollouts

$$\frac{S_{\text{model}} - S_{\text{random}}}{S_{\text{human}} - S_{\text{random}}}$$



10000 rollouts






# Human normalized scores

| Agent                   | Median            | Mean              | Env. Frames  |
|-------------------------|-------------------|-------------------|--------------|
| Ape-X                   | 434.1%            | 1695.6%           | 22.8B        |
| R2D2                    | 1920.6%           | 4024.9%           | <u>37.5B</u> |
| <i>MuZero</i>           | <b>2041.1%</b>    | <b>4999.2%</b>    | <u>20.0B</u> |
| IMPALA                  | 191.8%            | 957.6%            | 200M         |
| Rainbow                 | 231.1%            | —                 | 200M         |
| UNREAL                  | 250% <sup>a</sup> | 880% <sup>a</sup> | 250M         |
| LASER                   | 431%              | —                 | 200M         |
| <u>MuZero Reanalyze</u> | <b>731.1%</b>     | <b>2168.9%</b>    | 200M         |






A total  
20 TPU: 12h

400K


# References I

-  Aravind Rajeswaran, Igor Mordatch, and Vikash Kumar. “A game theoretic framework for model based reinforcement learning”. In: *arXiv preprint arXiv:2004.07804* (2020).
-  Ramanan Sekar et al. “Planning to Explore via Self-Supervised World Models”. In: *arXiv preprint arXiv:2005.05960* (2020).
-  Tianhe Yu et al. *MOPO: Model-based Offline Policy Optimization*. 2020. *arXiv: 2005.13239 [cs.LG]*.
-  David Ha and Jürgen Schmidhuber. “World models”. In: *arXiv preprint arXiv:1803.10122* (2018).
-  Lukasz Kaiser et al. “Model-based reinforcement learning for atari”. In: *arXiv preprint arXiv:1903.00374* (2019).

## References II

-  Aviv Tamar et al. “Value iteration networks”. In: *Advances in Neural Information Processing Systems*. 2016, pp. 2154–2162.
-  Junhyuk Oh, Satinder Singh, and Honglak Lee. “Value prediction network”. In: *Advances in Neural Information Processing Systems*. 2017, pp. 6118–6128.
-  David Silver et al. “The predictron: End-to-end learning and planning”. In: *Proceedings of the 34th International Conference on Machine Learning-Volume 70*. JMLR. org. 2017, pp. 3191–3199.
-  David Silver et al. “Mastering chess and shogi by self-play with a general reinforcement learning algorithm”. In: *arXiv preprint arXiv:1712.01815* (2017).
-  Julian Schrittwieser et al. “Mastering atari, go, chess and shogi by planning with a learned model”. In: *arXiv preprint arXiv:1911.08265* (2019).

## References III

-  Danijar Hafner et al. “Dream to control: Learning behaviors by latent imagination”. In: *arXiv preprint arXiv:1912.01603* (2019).